


# Multiple Linear Regression For Prediction



Source: Python plot posted to StackExchange <https://stackoverflow.com/questions/26431800/plot-linear-model-in-3d-with-matplotlib>

# Explanatory vs. Predictive

Explain/describe  
population relationships

- Small sample, few variables
- Retrospective
- Find good fitting regression model
- Confidence intervals, hypothesis test, p-value

Predict values of new  
records

- Large sample, many variables
- Prospective
- Regression with high predictive power
- Predictive power on holdout data

# Example – Used Toyota Car Prices



Scenario – Toyota dealers accept used cars in trade-ins, need predictive model to know how much to offer the customer

# The Regression Equation

The diagram illustrates the components of the regression equation  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \epsilon$ . Blue arrows point from descriptive labels to specific parts of the equation:

- outcome**: points to  $Y$
- predictors**: points to the  $x$  variables ( $x_1, x_2, \dots, x_p$ )
- constant**: points to  $\beta_0$
- coefficients**: points to the  $\beta$  parameters ( $\beta_1, \beta_2, \dots, \beta_p$ )
- Error (noise)**: points to  $\epsilon$

# The Regression Equation

Training phase – known outcome

predictors

Error (noise)

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \epsilon$$

constant

coefficients

The diagram shows the regression equation  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \epsilon$ . A red arrow points from the text 'Training phase – known outcome' to the variable  $Y$ . Blue arrows point from the word 'predictors' to the variables  $x_1$ ,  $x_2$ , and  $x_p$ . A blue arrow points from the text 'Error (noise)' to the error term  $\epsilon$ . A blue arrow points from the word 'constant' to the coefficient  $\beta_0$ . A blue arrow points from the word 'coefficients' to the coefficients  $\beta_1$ ,  $\beta_2$ , and  $\beta_p$ .

# The Regression Equation

Scoring  
phase -  
unknown  
outcome

The diagram shows the regression equation  $? = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \epsilon$ . Annotations include: a red arrow pointing to the red question mark from the text 'Scoring phase - unknown outcome'; blue arrows pointing from 'predictors' to  $x_1$ ,  $x_2$ , and  $x_p$ ; a blue arrow pointing from 'Constant (estimated in the training phase)' to  $\beta_0$ ; a blue arrow pointing from 'Coefficients (estimated in the training phase)' to  $\beta_1$ ,  $\beta_2$ , and  $\beta_p$ ; and a blue arrow pointing from 'Error (noise)' to  $\epsilon$ .

$$? = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \epsilon$$

predictors

Error (noise)

Constant (estimated in the training phase)

Coefficients (estimated in the training phase)

# Used Toyota Car Prices – Predictors



| Variable   | Description                              |
|------------|--|
| Price      | Offer price in Euros                     |
| Age        | Age in months as of August 2004          |
| Kilometers | Accumulated kilometers on odometer       |
| Fuel Type  | Fuel type ( <i>Petrol, Diesel, CNG</i> ) |
| HP         | Horsepower                               |
| Metallic   | Metallic color? (Yes = 1, No = 0)        |
| Automatic  | Automatic (Yes = 1, No = 0)              |
| CC         | Cylinder volume in cubic centimeters     |
| Doors      | Number of doors                          |
| QuartTax   | Quarterly road tax in Euros              |
| Weight     | Weight in kilograms                      |

# Import Needed Functionality

```
import itertools
import pandas as pd
import statsmodels.formula.api as sm
from sklearn.feature_selection import
    SelectKBest, f_regression
from utilities import regressionSummary
```



# Prepare and Partition Data

```
# Reduce data frame to the top 1000 rows and select
# columns for regression analysis
car_df = pd.read_csv(DATA / 'ToyotaCorolla.csv',
                     encoding='latin-1')
selected_var = [car_df.columns[i] for i in (2, 3, 6,
      7, 8, 9, 11, 12, 13, 16, 17)]
car_df = car_df.iloc[0:1000]

# Partition data
train_df = car_df.sample(frac=0.6, random_state=1)
valid_df = car_df.drop(train_df.index)
```

# Prepare and Partition Data

```
# Reduce data frame to the top 1000 rows and select
# columns for regression analysis
car_df = pd.read_csv(DATA / 'ToyotaCorolla.csv',
                     encoding='latin-1')
selected_var = [car_df.columns[i] for i in (2, 3, 6,
      7, 8, 9, 11, 12, 13, 16, 17)]
car_df = car_df.iloc[0:1000]

# Partition data
train_df = car_df.sample(frac=0.6, random_state=1)
valid_df = car_df.drop(train_df.index)
```

# Prepare and Partition Data

```
# Reduce data frame to the top 1000 rows and select
# columns for regression analysis
car_df = pd.read_csv(DATA / 'ToyotaCorolla.csv',
                     encoding='latin-1')
selected_var = [car_df.columns[i] for i in (2, 3, 6,
      7, 8, 9, 11, 12, 13, 16, 17)]
car_df = car_df.iloc[0:1000]

# Partition data
train_df = car_df.sample(frac=0.6, random_state=1)
valid_df = car_df.drop(train_df.index)
```

# Prepare and Partition Data

```
# Reduce data frame to the top 1000 rows and select
# columns for regression analysis
car_df = pd.read_csv(DATA / 'ToyotaCorolla.csv',
                     encoding='latin-1')
selected_var = [car_df.columns[i] for i in (2, 3, 6,
      7, 8, 9, 11, 12, 13, 16, 17)]
car_df = car_df.iloc[0:1000]

# Partition data
train_df = car_df.sample(frac=0.6, random_state=1)
valid_df = car_df.drop(train_df.index)
```

# Prepare and Partition Data

```
# Reduce data frame to the top 1000 rows and select
# columns for regression analysis
car_df = pd.read_csv(DATA / 'ToyotaCorolla.csv',
                     encoding='latin-1')
selected_var = [car_df.columns[i] for i in (2, 3, 6,
      7, 8, 9, 11, 12, 13, 16, 17)]
car_df = car_df.iloc[0:1000]

# Partition data
train_df = car_df.sample(frac=0.6, random_state=1)
valid_df = car_df.drop(train_df.index)
```

# Fit the regression

```
# Construct regression formula and use sm.ols to run a  
# linear regression of Price  
# on the remaining 11 predictors in the training set  
  
independent_var = list(selected_var)  
independent_var.remove('Price')  
formula = 'Price ~ ' + ' + '.join(independent_var)  
car_lm = sm.ols(formula=formula, data=train_df).fit()  
car_lm.summary()
```

# Fit the regression

```
# Construct regression formula and use sm.ols to run a  
# linear regression of Price  
# on the remaining 11 predictors in the training set  
  
independent_var = list(selected_var)  
independent_var.remove('Price')  
formula = 'Price ~ ' + ' + '.join(independent_var)  
car_lm = sm.ols(formula=formula, data=train_df).fit()  
car_lm.summary()
```

# Review Results

| OLS Regression Results |                  |                     |           |       |           |           |
|------------------------|------------------|---------------------|-----------|-------|-----------|-----------|
| Dep. Variable:         | Price            | R-squared:          | 0.867     |       |           |           |
| Model:                 | OLS              | Adj. R-squared:     | 0.864     |       |           |           |
| Method:                | Least Squares    | F-statistic:        | 347.4     |       |           |           |
| Date:                  | Mon, 18 Jun 2018 | Prob (F-statistic): | 7.53e-249 |       |           |           |
| Time:                  | 19:44:28         | Log-Likelihood:     | -5183.5   |       |           |           |
| No. Observations:      | 600              | AIC:                | 1.039e+04 |       |           |           |
| Df Residuals:          | 588              | BIC:                | 1.044e+04 |       |           |           |
| Df Model:              | 11               |                     |           |       |           |           |
| Covariance Type:       | nonrobust        |                     |           |       |           |           |
|                        | coef             | std err             | t         | P> t  | [0.025    | 0.975]    |
| Intercept              | -6314.1254       | 1822.923            | -3.464    | 0.001 | -9894.358 | -2733.893 |
| Age_08_04              | -127.4934        | 4.896               | -26.038   | 0.000 | -137.110  | -117.877  |
| KM                     | -0.0206          | 0.002               | -8.929    | 0.000 | -0.025    | -0.016    |
| HP                     | 36.8883          | 5.115               | 7.212     | 0.000 | 26.843    | 46.933    |
| Met_Color              | 82.3558          | 123.302             | 0.668     | 0.504 | -159.810  | 324.521   |
| Automatic              | 143.5263         | 289.166             | 0.496     | 0.620 | -424.397  | 711.449   |
| CC                     | 0.0236           | 0.097               | 0.243     | 0.808 | -0.167    | 0.215     |
| Doors                  | -46.6678         | 62.884              | -0.742    | 0.458 | -170.172  | 76.836    |
| Quarterly_Tax          | 12.1313          | 2.851               | 4.255     | 0.000 | 6.532     | 17.730    |
| Weight                 | 18.1393          | 1.822               | 9.954     | 0.000 | 14.560    | 21.718    |
| Fuel_Type_Diesel       | -31.4359         | 569.873             | -0.055    | 0.956 | -1150.670 | 1087.799  |
| Fuel_Type_Petrol       | 1461.2586        | 574.021             | 2.546     | 0.011 | 333.877   | 2588.640  |
| Omnibus:               | 112.122          | Durbin-Watson:      | 1.946     |       |           |           |
| Prob(Omnibus):         | 0.000            | Jarque-Bera (JB):   | 1305.165  |       |           |           |
| Skew:                  | -0.420           | Prob(JB):           | 3.86e-284 |       |           |           |
| Kurtosis:              | 10.176           | Cond. No.           | 2.34e+06  |       |           |           |



# Review Results

| OLS Regression Results |                  |                     |           |       |           |           |
|------------------------|------------------|---------------------|-----------|-------|-----------|-----------|
| Dep. Variable:         | Price            | R-squared:          | 0.867     |       |           |           |
| Model:                 | OLS              | Adj. R-squared:     | 0.864     |       |           |           |
| Method:                | Least Squares    | F-statistic:        | 347.4     |       |           |           |
| Date:                  | Mon, 18 Jun 2018 | Prob (F-statistic): | 7.53e-249 |       |           |           |
| Time:                  | 19:44:28         | Log-Likelihood:     | -5183.5   |       |           |           |
| No. Observations:      | 600              | AIC:                | 1.039e+04 |       |           |           |
| Df Residuals:          | 588              | BIC:                | 1.044e+04 |       |           |           |
| Df Model:              | 11               |                     |           |       |           |           |
| Covariance Type:       | nonrobust        |                     |           |       |           |           |
|                        | coef             | std err             | t         | P> t  | [0.025    | 0.975]    |
| Intercept              | -6314.1254       | 1822.923            | -3.464    | 0.001 | -9894.358 | -2733.893 |
| Age_08_04              | -127.4934        | 4.896               | -26.038   | 0.000 | -137.110  | -117.877  |
| KM                     | -0.0206          | 0.002               | -8.929    | 0.000 | -0.025    | -0.016    |
| HP                     | 36.8883          | 5.115               | 7.212     | 0.000 | 26.843    | 46.933    |
| Met_Color              | 82.3558          | 123.302             | 0.668     | 0.504 | -159.810  | 324.521   |
| Automatic              | 143.5263         | 289.166             | 0.496     | 0.620 | -424.397  | 711.449   |
| CC                     | 0.0236           | 0.097               | 0.243     | 0.808 | -0.167    | 0.215     |
| Doors                  | -46.6678         | 62.384              | -0.742    | 0.458 | -170.172  | 76.836    |
| Quarterly_Tax          | 12.1313          | 2.851               | 4.255     | 0.000 | 6.532     | 17.730    |
| Weight                 | 18.1393          | 3.822               | 9.954     | 0.000 | 14.560    | 21.718    |
| Fuel_Type_Diesel       | -31.4359         | 509.873             | -0.055    | 0.956 | -1150.670 | 1087.799  |
| Fuel_Type_Petrol       | 1461.2586        | 574.021             | 2.546     | 0.011 | 333.877   | 2588.640  |
| Omnibus:               | 112.122          | Durbin-Watson:      | 1.946     |       |           |           |
| Prob(Omnibus):         | 0.000            | Jarque-Bera (JB):   | 1305.165  |       |           |           |
| Skew:                  | -0.420           | Prob(JB):           | 3.86e-284 |       |           |           |
| Kurtosis:              | 10.176           | Cond. No.           | 2.34e+06  |       |           |           |

# Score the Validation Data

```
# Use predict() to make predictions on a new set

car_lm_pred = car_lm.predict(valid_df)
result = pd.DataFrame({'Predicted': car_lm_pred,
                       'Actual': valid_df.Price, 'Residual':
                       valid_df.Price - car_lm_pred })
```

# Look at First 10 Rows - Validation

Output

|    | Actual | Predicted    | Residual     |
|----|--------|--------------|--------------|
| 1  | 13750  | 16206.480671 | -2456.480671 |
| 7  | 18600  | 16704.259643 | 1895.740357  |
| 10 | 20950  | 21003.848119 | -53.848119   |
| 15 | 22000  | 20883.866334 | 1116.133666  |
| 20 | 15950  | 15039.157810 | 910.842190   |
| 21 | 16950  | 17078.992106 | -128.992106  |
| 22 | 15950  | 15799.618511 | 150.381489   |
| 24 | 16250  | 16302.187336 | -52.187336   |
| 25 | 15950  | 16757.750599 | -807.750599  |
| 26 | 17495  | 16377.003081 | 1117.996919  |

# Mean Error (ME)

Output

| Actual |       | Predicted    | Residual     | Average<br>of the<br>residuals |
|--------|-------|--------------|--------------|--------------------------------|
| 1      | 13750 | 16206.480671 | -2456.480671 |                                |
| 7      | 18600 | 16704.259643 | 1895.740357  |                                |
| 10     | 20950 | 21003.848119 | -53.848119   |                                |
| 15     | 22000 | 20883.866334 | 1116.133666  |                                |
| 20     | 15950 | 15039.157810 | 910.842190   |                                |
| 21     | 16950 | 17078.992106 | -128.992106  |                                |
| 22     | 15950 | 15799.618511 | 150.381489   |                                |
| 24     | 16250 | 16302.187336 | -52.187336   |                                |
| 25     | 15950 | 16757.750599 | -807.750599  |                                |
| 26     | 17495 | 16377.003081 | 1117.996919  |                                |

# Mean Error

Output

| Actual   | Predicted    | Residual     | Average<br>\$280 |
|----------|--------------|--------------|------------------|
| 1 13750  | 16206.480671 | -2456.480671 |                  |
| 7 18600  | 16704.259643 | 1895.740357  |                  |
| 10 20950 | 21003.848119 | -53.848119   |                  |
| 15 22000 | 20883.866334 | 1116.133666  |                  |
| 20 15950 | 15039.157810 | 910.842190   |                  |
| 21 16950 | 17078.992106 | -128.992106  |                  |
| 22 15950 | 15799.618511 | 150.381489   |                  |
| 24 16250 | 16302.187336 | -52.187336   |                  |
| 25 15950 | 16757.750599 | -807.750599  |                  |
| 26 17495 | 16377.003081 | 1117.996919  |                  |

# Root Mean Squared Error (RMSE)

Output

| Actual |       | Predicted    | Residual     | Square the residuals, average them, take square root |
|--------|-------|--------------|--------------|--|
| 1      | 13750 | 16206.480671 | -2456.480671 |  |
| 7      | 18600 | 16704.259643 | 1895.740357  |  |
| 10     | 20950 | 21003.848119 | -53.848119   |  |
| 15     | 22000 | 20883.866334 | 1116.133666  |  |
| 20     | 15950 | 15039.157810 | 910.842190   |  |
| 21     | 16950 | 17078.992106 | -128.992106  |  |
| 22     | 15950 | 15799.618511 | 150.381489   |  |
| 24     | 16250 | 16302.187336 | -52.187336   |  |
| 25     | 15950 | 16757.750599 | -807.750599  |  |
| 26     | 17495 | 16377.003081 | 1117.996919  |  |

# Mean Absolute Error (MAE)

Output

|    | Actual | Predicted    | Residual     |
|----|--------|--------------|--------------|
| 1  | 13750  | 16206.480671 | -2456.480671 |
| 7  | 18600  | 16704.259643 | 1895.740357  |
| 10 | 20950  | 21003.848119 | -53.848119   |
| 15 | 22000  | 20883.866334 | 1116.133666  |
| 20 | 15950  | 15039.157810 | 910.842190   |
| 21 | 16950  | 17078.992106 | -128.992106  |
| 22 | 15950  | 15799.618511 | 150.381489   |
| 24 | 16250  | 16302.187336 | -52.187336   |
| 25 | 15950  | 16757.750599 | -807.750599  |
| 26 | 17495  | 16377.003081 | 1117.996919  |

Take absolute  
values of  
residuals, find  
average

# Mean Absolute Percentage Error (MAPE)

Output

|    | Actual | Predicted    | Residual     |
|----|--------|--------------|--------------|
| 1  | 13750  | 16206.480671 | -2456.480671 |
| 7  | 18600  | 16704.259643 | 1895.740357  |
| 10 | 20950  | 21003.848119 | -53.848119   |
| 15 | 22000  | 20883.866334 | 1116.133666  |
| 20 | 15950  | 15039.157810 | 910.842190   |
| 21 | 16950  | 17078.992106 | -128.992106  |
| 22 | 15950  | 15799.618511 | 150.381489   |
| 24 | 16250  | 16302.187336 | -52.187336   |
| 25 | 15950  | 16757.750599 | -807.750599  |
| 26 | 17495  | 16377.003081 | 1117.996919  |

Take residual as  
% of actual, find  
absolute value,  
find average



# Review Error Metrics - Validation

## Regression statistics

|                                       |             |
|---------------------------------------|-------------|
| Mean Error (ME)                       | : 49.1865   |
| Root Mean Squared Error (RMSE)        | : 1349.2010 |
| Mean Absolute Error (MAE)             | : 1021.6062 |
| Mean Percentage Error (MPE)           | : 0.0381    |
| Mean Absolute Percentage Error (MAPE) | : 9.0735    |

# Used Toyota Car Prices – Predictors

| Variable   | Description                              |
|------------|--|
| Price      | Offer price in Euros                     |
| Age        | Age in months as of August 2004          |
| Kilometers | Accumulated kilometers on odometer       |
| Fuel Type  | Fuel type ( <i>Petrol, Diesel, CNG</i> ) |
| HP         | Horsepower                               |
| Metallic   | Metallic color? (Yes = 1, No = 0)        |
| Automatic  | Automatic (Yes = 1, No = 0)              |
| CC         | Cylinder volume in cubic centimeters     |
| Doors      | Number of doors                          |
| QuartTax   | Quarterly road tax in Euros              |
| Weight     | Weight in kilograms                      |